



From expert responses to a shared claim set

Round 1 responses → Extracted claims → Grouped claims for round 2

a Before → after grouping



b How much meaning survives?



c Which information is lost?



Figure 1 | Preserving meaning from expert responses to grouped claims.

a, Coverage and faithfulness before grouping (purple) and after grouping (orange), shown for the same 320 simulated consultations. Coverage measures how much original claim content is retained; faithfulness measures whether the resulting claims preserve the source meaning. Both maps have equal x-y scaling, identical 0–100% axes and common density normalisation; the corner diamond marks complete, faithful retention. Contours enclose 50% and 90% of each estimated density; marginal curves show each metric separately. The grouping changes are constructed illustrations, not measured effects. b, Saved assistant-pilot retention across four synthetic scenarios (six panels and 55 reference-claim occurrences per scenario). Purple marks full preservation, amber adds partial preservation, and the pale remainder represents omissions. Judgments were unblinded self-review; all omissions reflect the eight-claim cap. c, Illustrative retention distributions for four information types, using a separate 120 simulated cases per type. Faint points show individual cases; prominent dots and horizontal segments show medians and interquartile ranges, not confidence intervals. These panels do not establish a measured grouping effect or independently validated platform performance.