


From expert responses to a shared claim set

Round 1 responses → Extracted claims → Grouped claims for round 2


Extraction and grouping

MEASURED · MODEL-JUDGED SYNTHETIC INPUTS

A Before grouping




B After grouping



C How much meaning survives?

AFTER GROUPING · FULL AND PARTIAL RETENTION



D Which information is lost?

AFTER GROUPING · INFORMATION TYPES



Figure 1 | Preserving meaning from expert responses to grouped claims.

A–B, Exact coverage and faithfulness scores from 76 paired synthetic panels; purple is extraction and copper is grouping. Marginal ticks show empirical frequencies. Coincident panels are counted, not displaced; density contours are used only when at least 30 panels and six distinct coordinates are available. C, Full and partial retention after grouping, pooled over reference occurrences in the available nested sizes for each topic. D, Individual panel retention by reference category; dots and lines show medians and interquartile ranges, not confidence intervals. Vertical jitter separates observations without changing scores. Density estimates, where shown, use reflected Gaussian smoothing with a fixed 3.5-percentage-point bandwidth. Sources and references were model-authored and frozen before extraction; judgments are provisional and not human validated. Similar meanings from different speakers may be counted separately. Grouping is capped at eight statements. This is a benchmark of prompts, not the deployed platform pipeline; no causal size effect is established.