


From expert responses to a shared claim set

36 completed paired panels · six topics · 4, 8 and 16 participants · real OpenRouter calls

Purple: extracted claims · Copper: grouped claims · Dots: completed panels; larger markers count overlap



Controlled synthetic benchmark · provisional model-judged results

Coverage is the fraction of unique source propositions fully preserved; partial preservation does not receive full credit. Faithfulness is the fraction of output claims judged fully supported. Sizes are nested within topic and replicate. Panels c-d show individual panels and unweighted means, not confidence intervals. A separate model judges each stage against original source text. These are benchmark prompts, not a test of the deployed Symphonia extraction pipeline. Templated inputs and the eight-claim grouping cap limit generalisation. Full inputs, outputs and audits are downloadable.