


CLAIM EXTRACTION

How much survives—and how faithfully?

DESIGN PREVIEW · 320 simulated points · Not measured Symphonia performance



One dot represents one hypothetical consultation. Faint dots show individual variation; shading and contours show concentration; marginal curves show each axis separately.


Contours enclose 50% and 90% of the smoothed illustrative density, not confidence intervals. All coordinates are generated solely to preview the figure design. No models were evaluated.

Whose evidence survives?


ILLUSTRATIVE ONLY · Simulated data · No measured platform performance

A distribution reveals the typical case. Paired perturbations reveal selective treatment.

a Retention by information type



b A matched bias test




Top: 120 hypothetical cases per category. Bottom: 120 hypothetical label-swap pairs.
Fine points retain individual cases; bold marks show medians and interquartile ranges.

DESIGN SIMULATION — NOT RESULTS. Curves are descriptive KDEs, not confidence intervals. Category differences alone do not establish bias. The matched test changes only the status label; its effect distribution here is invented for design purposes.

Does the summary change the evidence?

ILLUSTRATIVE ONLY · Simulated data · No measured platform performance

An accurate summary sits on the diagonal. Systematic departures reveal distortion.



Each point is one hypothetical rated claim (240 total). The diagonal means exact reporting. The right panel separates error direction and spread, so a mean near zero cannot hide large mistakes.


DESIGN SIMULATION — NOT RESULTS. Contours show 50% / 90% of smoothed illustrative density, not statistical uncertainty. Agreement percentages and error patterns are simulated; no actual consultation ratings enter this figure.

Agreement is not the same as accuracy


ILLUSTRATIVE ONLY · Simulated data · No measured platform performance

Separate convergence from correction—and compare against a no-feedback condition.

a Within-panel trajectories



b What kind of convergence?



Left: 100 hypothetical paired panels per condition; bold lines show arithmetic means.

Right: peer-feedback panels can converge while becoming either more or less accurate.


DESIGN SIMULATION — NOT RESULTS. Disagreement and factual-accuracy scores are constructed here, not calculated from recorded ballots. Actual accuracy requires decidable claims; preferences have no truth score. Contours describe simulated density.

Does information loss change decisions?


ILLUSTRATIVE ONLY · Simulated data · No measured platform performance

The final test is what readers understand and decide—not how persuasive a summary sounds.


a Consequences across cases



b Paired difference by case



c Is confidence warranted?



Hypothetical reader-study design: 120 cases compared across three evidence presentations.
Paired differences isolate the comparison; calibration checks whether confidence is deserved.

DESIGN SIMULATION — NOT RESULTS. Case-level losses and calibration curves are entirely constructed. Shifts are design examples, not predicted benefits. A real study needs randomized readers, a prespecified decision rule, and reader- and case-aware analysis.